

```
# ----- 1) IMPORTS & DATA CLEANING -----
import pandas as pd
import string, re, nltk
from nltk.corpus import stopwords

# Download NLTK stop-word list (one-time per environment)
nltk.download('stopwords')

# Load the CSV that contains the raw text and sentiment labels
df = pd.read_csv("legal_sentiment_dataset.csv") # Columns: Text, Sentiment

# Create a reusable cleaning function
stop_words = set(stopwords.words('english'))

def clean_text(text: str) -> str:
    text = text.lower() # a) lowercase
    text = re.sub(r"\d+", "", text) # b) remove numbers
    text = text.translate(str.maketrans("", "", string.punctuation)) # c) strip punctuation
    text = " ".join(w for w in text.split() if w not in stop_words) # d) remove stop-words
    return text

# Apply cleaning and store in new column
df["clean_text"] = df["Text"].apply(clean_text)

# Show a small sample
print(df.head(3))
```



```
Text Sentiment \
0 [Case 3708] The appeal is hereby dismissed for... Negative
1 [Case 6994] The plaintiff's motion for summary... Positive
2 [Case 2793] The legal claim is barred by the s... Negative
```

```
clean_text
0 case appeal hereby dismissed lack merit
1 case plaintiffs motion summary judgment granted
2 case legal claim barred statute limitations
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

Pl
fo
or
[bl](#)
to
se
m
in
al
ne
fe
tip
ar
tr
ar
fe
no
st
as
[A](#)
[a](#)
[B](#)
[E](#)
[w](#)
[C](#)

2
0
1

```
# ----- 2) TRAIN-TEST SPLIT & TF-IDF -----
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer

# Separate features & labels
X = df["clean_text"]
y = df["Sentiment"]

# 80 % training / 20 % test
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42, stratify=y
)

# TF-IDF converts text to numeric vectors
vectorizer = TfidfVectorizer(max_features=5000) # top 5 000 terms
X_train_vec = vectorizer.fit_transform(X_train) # fit + transform (train)
X_test_vec = vectorizer.transform(X_test) # transform only (test)

print("Training matrix shape:", X_train_vec.shape)
print("Example feature names:", vectorizer.get_feature_names_out()[:10])

⇒ Training matrix shape: (80, 72)
Example feature names: ['acquitted' 'agreement' 'amicable' 'appeal' 'barred' 'beyond' 'breach'
'case' 'caused' 'charges']

# ----- 3) MODEL, METRICS, PREDICTION, SAVE -----
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
import pickle

# a) Train a binary classifier
model = LogisticRegression(max_iter=1000) # more iterations for convergence
model.fit(X_train_vec, y_train)

# b) Evaluate on the held-out test set
y_pred = model.predict(X_test_vec)
```

```

print("Accuracy on test set:", accuracy_score(y_test, y_pred))
print("\nDetailed classification report:\n", classification_report(y_test, y_pred))

# c) Helper for predictions on new sentences
def predict_sentiment(raw_text: str) -> str:
    cleaned = clean_text(raw_text)
    vec = vectorizer.transform([cleaned])
    return model.predict(vec)[0]

print("\nSample inference:",
      predict_sentiment("The injunction is granted to prevent further harm."))

# d) Persist model + vectorizer for later reuse
with open("legal_sentiment_model.pkl", "wb") as f:
    pickle.dump((model, vectorizer), f)
print("\nModel saved as legal_sentiment_model.pkl")

```

➡ Accuracy on test set: 1.0

Detailed classification report:

	precision	recall	f1-score	support
Negative	1.00	1.00	1.00	9
Positive	1.00	1.00	1.00	11
accuracy			1.00	20
macro avg	1.00	1.00	1.00	20
weighted avg	1.00	1.00	1.00	20

Sample inference: Positive

Model saved as legal_sentiment_model.pkl

in

2
0
1

201

202

200

2
1
1

2
1
0

200

201.

2
1
0

2
0
3

2
0
2

2
0
2

200

2
0
2